

Box 1 - Example DFA workflow

Model specification

For this example we will assume that the observations at time t (\mathbf{y}_t) arise from a multivariate normal distribution wherein the mean vector is the sum of a vector of intercepts ($\boldsymbol{\alpha}$) and latent factors ($\mathbf{L}\boldsymbol{\gamma}_t$), and the covariance matrix $\boldsymbol{\Phi}$ is diagonal and unequal. Eqn (1) then becomes:

$$\mathbf{y}_t \sim \text{MVN}(\boldsymbol{\alpha} + \mathbf{L}\boldsymbol{\gamma}_t, \boldsymbol{\Phi}).$$

The latent factors are themselves a multivariate random walk with normal errors. Here we assume \mathbf{Q} equals the identity matrix \mathbf{I} , such that Eqn (3) is

$$\boldsymbol{\gamma}_t \sim \text{MVN}(\boldsymbol{\gamma}_{t-1}, \mathbf{Q}).$$

Simulated data

Here are some simulated data for 5 time series with 30 measurements each based upon 2 latent factors. The observation variances are also unequal. We will set the covariance of the process errors \mathbf{Q} to the identity matrix \mathbf{I} . Note that we will also subtract the mean of the data so as to aid in parameter estimation (see below).

```
set.seed(123)
## matrix dims
nn <- 5; mm <- 2; tt <- 30
## zero-mean latent factors
gamma1 <- cumsum(rnorm(tt)); gamma1 <- gamma1 - mean(gamma1)
gamma2 <- cumsum(rnorm(tt)); gamma2 <- gamma2 - mean(gamma2)
## factor loadings
LL <- matrix(c(0.2, 0.9,
               0.4, 0.7,
               0.6, 0.5,
               0.8, 0.3,
               1.0, 0.1),
             nn, mm, byrow=TRUE)
## covariance matrix for observation errors
Phi <- diag(seq(1,5)/5)
## observation errors
vv <- t(MASS::mvrnorm(tt, matrix(0,nn,1), Phi))
## simulated data
yy <- LL %*% rbind(gamma1,gamma2) + vv
## zero-mean data
yy <- t(scale(t(yy), scale=FALSE))
```

Fitting the model

There are several options for fitting factor models, but for this example we make use of the **MARSS** package (Holmes et al. 2012) for **R** to estimate the parameters in the DFA model. The notation of

Holmes et al. is somewhat different than ours, however, so we must write out our DFA model in a form based on the following specification of a state-space model:

$$\begin{aligned}\mathbf{y}_t &\sim \text{MVN}(\mathbf{a} + \mathbf{Z}\mathbf{x}_t, \mathbf{R}) \\ \mathbf{x}_t &\sim \text{MVN}(\mathbf{u} + \mathbf{B}\mathbf{x}_{t-1}, \mathbf{Q}).\end{aligned}$$

Thus, for our DFA model we want $\mathbf{a} = \boldsymbol{\alpha}$, $\mathbf{Z} = \mathbf{L}$, $\mathbf{x}_t = \boldsymbol{\gamma}_t$, $\mathbf{R} = \boldsymbol{\Phi}$, $\mathbf{u} = \mathbf{0}$, and $\mathbf{B} = \mathbf{I}$. Furthermore, because we subtracted the mean of each time series, we will set $\mathbf{a} = \mathbf{0}$. Note that this will not affect our estimates of the factors and their loadings.

The `MARSS` function makes use of list matrices, which allow one to mix character and numeric classes within the same matrix. Any entries of class `character` are interpreted as parameters to be estimated; any `numeric` entries are fixed at their value.

```
library(MARSS)
## model list
mod_list <- list(
  ## observation eqn
  A = matrix(0,nn,1),          # zero vector
  Z = matrix(list(0),nn,mm),    # all 0's for now
  R = matrix(list(0),nn,nn),    # all 0's for now
  ## state eqn
  U = matrix(0,mm,1),          # zero vector
  B = diag(mm),                # identity matrix
  Q = diag(mm)                 # identity matrix
)
## specify observation parameters
## upper right corner of Z stays zero
for(cc in 1:mm) {
  for(rr in cc:nn) {
    mod_list$Z[rr,cc] <- paste0(rr,cc)
  }
}
## R is diagonal and unequal
diag(mod_list$R) <- paste0(seq(nn),seq(nn))
## fit the model
dfa <- MARSS(yy, model=mod_list, inits=list(x0=matrix(0,2,1)))
```

Factor rotations

For this example we will use a varimax rotation of the factors and loadings, where we seek an $m \times m$ non-singular rotation matrix \mathbf{H} , such that the following state-space forms of our DFA model are equivalent (\mathbf{v}_t and \mathbf{w}_t are vectors of observation and process errors, respectively):

$$\begin{aligned}\mathbf{y}_t &= \mathbf{L}\boldsymbol{\gamma}_t + \boldsymbol{\alpha} + \mathbf{v}_t \Leftrightarrow \mathbf{y}_t = \mathbf{LH}^{-1}\boldsymbol{\gamma}_t + \boldsymbol{\alpha} + \mathbf{v}_t \\ \boldsymbol{\gamma}_t &= \boldsymbol{\gamma}_{t-1} + \mathbf{w}_t \Leftrightarrow \mathbf{H}\boldsymbol{\gamma}_t = \mathbf{H}\boldsymbol{\gamma}_{t-1} + \mathbf{H}\mathbf{w}_t.\end{aligned}$$

This is easily done in **R** via the `varimax` function:


```
## loadings matrix
L_hat <- coef(dfa, type="matrix")$Z
```

```
## inverse of the rotation matrix
H_inv <- varimax(L_hat)$rotmat
## rotate loadings matrix
L_rot <- L_hat %*% H_inv
## rotate factors
gamma_rot <- solve(H_inv) %*% dfa$states
## model fits
yhat <- L_rot %*% gamma_rot
```

Estimated parameters

Here are plots of the true (solid) and estimated factors (dashed). Note that **MARSS** has no way of ordering the factors, so care should be used when examining them. In this case, the estimated factors were interchanged.

```
matplot(cbind(gamma2,gamma1,t(gamma_rot)), type="l", lty=c(1,1,2,2), lwd=2,
        col=c("darkblue", "darkred"), ylab="Factors", xlab="Time")
```



Here is a comparison of the true factor loadings and the estimates in $\hat{\mathbf{L}}$.

```
## true values
LL
##      [,1] [,2]
## [1,] 0.2  0.9
## [2,] 0.4  0.7
## [3,] 0.6  0.5
## [4,] 0.8  0.3
## [5,] 1.0  0.1

## estimates
round(L_rot[,2:1],2)
##      [,1] [,2]
## [1,] 0.27 0.86
## [2,] 0.44 0.63
## [3,] 0.64 0.47
## [4,] 0.78 0.28
```

```
## [5,] 0.91 0.34
```

It looks like the estimated loadings are generally pretty close to the true values, with the exception of the loading of factor 2 on data series 5.

Model fits

Here is a plot of the simulated data (numbers) and the fits (lines) from the estimated DFA model.

```
matplot(t(yy), type="p", cex=0.8, col=rainbow(nn,start=0.55, end=0.85),
        ylab="Observations", xlab="Time")
matlines(t(yhat), type="l", lty=1, lwd=2, col=rainbow(nn,start=0.55, end=0.85))
```

